

Towards Efficient Verification of Parallel Applications with Mc SimGrid


Joint work with Martin Quinson (Magellan) and
Thierry Jéron (Devine)

Mathieu Laurent

February 20, 2025



Distributed computing



- HPC applications are distributed and concurrent
- Data shared via messages (e.g. MPI) or synchronizations (e.g. thread)
- Causes non-deterministic bugs
- Software model checking covers all cases

Content of this talk


- 1 Introduction
- 2 Dynamic software model checking
 - Principle
 - Partial order reduction
 - Best First (O)DPOR
 - Best First (O)DPOR
- 3 Explainability
- 4 Conclusion

(with ongoing work 📢)

Exploring the transition systems

A small MPI example


| | |
|---------------|--------|
| P_1/P_2 | P_3 |
| Send(P_3) | Recv() |
| | Recv() |



Exploring the transition systems

A small MPI example


| | |
|---------------|--------|
| P_1/P_2 | P_3 |
| Send(P_3) | Recv() |
| | Recv() |



Exploring the transition systems

A small MPI example


| | |
|---------------|--------|
| P_1/P_2 | P_3 |
| Send(P_3) | Recv() |
| | Recv() |



Exploring the transition systems

A small MPI example


| | |
|---------------|--------|
| P_1/P_2 | P_3 |
| Send(P_3) | Recv() |
| | Recv() |



Exploring the transition systems

A small MPI example


| | |
|---------------|--------|
| P_1/P_2 | P_3 |
| Send(P_3) | Recv() |
| | Recv() |



Exploring the transition systems

A small MPI example


| | |
|---------------|--------|
| P_1/P_2 | P_3 |
| Send(P_3) | Recv() |
| | Recv() |



Exploring the transition systems

A small MPI example


| | |
|---------------|--------|
| P_1/P_2 | P_3 |
| Send(P_3) | Recv() |
| | Recv() |



Exploring the transition systems

A small MPI example


| | |
|---------------|--------|
| P_1/P_2 | P_3 |
| Send(P_3) | Recv() |
| | Recv() |



Exploring the transition systems

A small MPI example


| | |
|---------------|--------|
| P_1/P_2 | P_3 |
| Send(P_3) | Recv() |
| | Recv() |



Exploring the transition systems

A small MPI example


| | |
|---------------|--------|
| P_1/P_2 | P_3 |
| Send(P_3) | Recv() |
| | Recv() |



Exploring the transition systems

A small MPI example

| | |
|---------------|--------|
| P_1/P_2 | P_3 |
| Send(P_3) | Recv() |
| | Recv() |




Exploring the transition systems

A small MPI example


| | |
|---------------|--------|
| P_1/P_2 | P_3 |
| Send(P_3) | Recv() |
| | Recv() |

Stateful exploration

15 states for 2 behaviors.



Stateless model checking





Stateless exploration

35 states for the same 2 behaviors.

Transition dependency

Two actions a_1, a_2 are **independent** if:




Example of two adjacent independent actions

Mazurkiewicz's traces [Maz'77]


Equivalence class of executions with adjacent independent actions swapped

DPOR approach [Fla'05]



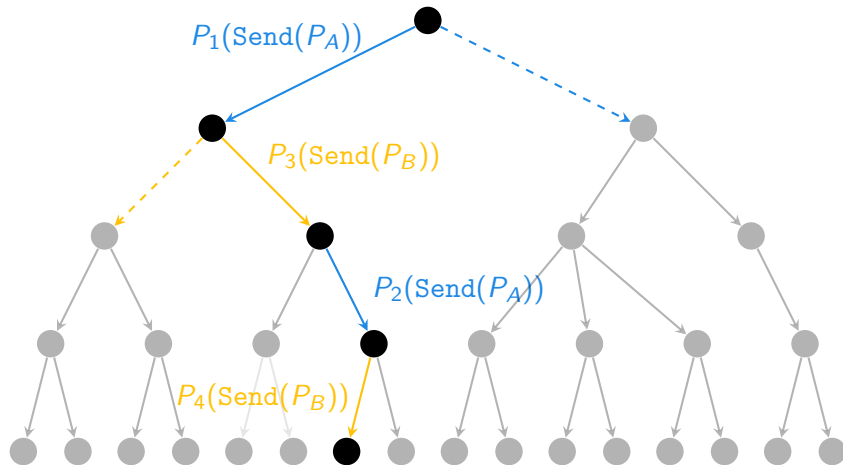
Classical depth first search algorithm

DPOR approach [Fla'05]




Start with an arbitrary execution

DPOR approach [Fla'05]




Discover dependencies

DPOR approach [Fla'05]




Recursive DFS exploration of what has been added

DPOR: persistent sets




A set T of transitions from s is **persistent** iff for any **sequence of transitions not in T** starting from s , t' is independent with T

It is sufficient to only explore transitions in a **persistent** set:



equivalent
to



DPOR: persistent sets

Five transitions enabled in s

$P_1(\text{Send}(P_A))$

$P_2(\text{Send}(P_A))$

$P_3(\text{Send}(P_B))$

$P_4(\text{Send}(P_B))$

$P_5(\text{LocalOp}())$

- $\{P_1, P_2\}, \{P_3, P_4, P_5\}, \{P_1, P_2, P_3, P_4\}, \dots$ are **persistent** sets
- $\{P_1\}, \{P_2, P_3\}, \dots$ are not persistent sets


DPOR: persistent sets

Five transitions enabled in s

$P_1(\text{Send}(P_A))$
 $P_2(\text{Send}(P_A))$
 $P_3(\text{Send}(P_B))$
 $P_4(\text{Send}(P_B))$
 $P_5(\text{LocalOp}())$

- $\{P_1, P_2\}, \{P_3, P_4, P_5\}, \{P_1, P_2, P_3, P_4\}, \dots$ are **persistent** sets
- $\{P_1\}, \{P_2, P_3\}, \dots$ are not persistent sets

DPOR builds persistent sets iteratively for each state



Pick an arbitrary transition

DPOR: persistent sets

Five transitions enabled in s

$P_1(\text{Send}(P_A))$

$P_2(\text{Send}(P_A))$


$P_3(\text{Send}(P_B))$

$P_4(\text{Send}(P_B))$

$P_5(\text{LocalOp}())$

- $\{P_1, P_2\}, \{P_3, P_4, P_5\}, \{P_1, P_2, P_3, P_4\}, \dots$ are **persistent** sets
- $\{P_1\}, \{P_2, P_3\}, \dots$ are not persistent sets

DPOR builds persistent sets iteratively for each state



Explore the corresponding subtree

DPOR: persistent sets

Five transitions enabled in s

$P_1(\text{Send}(P_A))$

$P_2(\text{Send}(P_A))$


$P_3(\text{Send}(P_B))$

$P_4(\text{Send}(P_B))$

$P_5(\text{LocalOp}())$

- $\{P_1, P_2\}, \{P_3, P_4, P_5\}, \{P_1, P_2, P_3, P_4\}, \dots$ are **persistent** sets
- $\{P_1\}, \{P_2, P_3\}, \dots$ are not persistent sets

DPOR builds persistent sets iteratively for each state



If a dependent transition is found, add it to the **persistent** set


DPOR: persistent sets

Five transitions enabled in s

$P_1(\text{Send}(P_A))$
 $P_2(\text{Send}(P_A))$
 $P_3(\text{Send}(P_B))$
 $P_4(\text{Send}(P_B))$
 $P_5(\text{LocalOp}())$

- $\{P_1, P_2\}, \{P_3, P_4, P_5\}, \{P_1, P_2, P_3, P_4\}, \dots$ are **persistent** sets
- $\{P_1\}, \{P_2, P_3\}, \dots$ are not persistent sets

DPOR builds persistent sets iteratively for each state



Repeat until no more dependent transition not in the **set** is found

DPOR: sleep sets

Sleep set

For each prefix E

- $sleep(E)$ contains the transitions already explored from E
- $sleep(E \cdot t)$ is initialized with $\{t' \mid t' \in sleep(E) \text{ and } t' \text{ is independent with } t\}$


DPOR: sleep sets

Sleep set

For each prefix E

- $\text{sleep}(E)$ contains the transitions already explored from E
- $\text{sleep}(E \cdot t)$ is initialized with $\{t' \mid t' \in \text{sleep}(E) \text{ and } t' \text{ is independent with } t\}$

It is sound to never explore a transition in a **sleep** set:



After exploring the subtree starting with t , $\text{sleep}(E) = \{t\}$


DPOR: sleep sets

Sleep set

For each prefix E

- $\text{sleep}(E)$ contains the transitions already explored from E
- $\text{sleep}(E \cdot t)$ is initialized with $\{t' \mid t' \in \text{sleep}(E) \text{ and } t' \text{ is independent with } t\}$

It is sound to never explore a transition in a **sleep** set:



At any time when exploring t' , if t is still in the **sleep** set at u ...


DPOR: sleep sets

Sleep set

For each prefix E


- $\text{sleep}(E)$ contains the transitions already explored from E
- $\text{sleep}(E \cdot t)$ is initialized with $\{t' \mid t' \in \text{sleep}(E) \text{ and } t' \text{ is independent with } t\}$

It is sound to never explore a transition in a **sleep** set:



for each execution starting by t from u , an **equivalent** has been explored

ODPOR approach [Abd'14]




Compute initials and handle a tree of sequences instead of a single step

ODPOR: better sets

Weak Initials

$p \in WI_{[E]}(w)$ iff there exists v, w' such that



To explore w from E , we can start by any process in $WI_{[E]}(w)$


ODPOR: better sets

Weak Initials


$p \in WI_{[E]}(w)$ iff there exists v, w' such that



is equivalent to



To explore w from E , we can start by any process in $WI_{[E]}(w)$



Source sets computation


ODPOR: better sets

Weak Initials


$p \in WI_{[E]}(w)$ iff there exists v, w' such that



is equivalent to



To explore w from E , we can start by any process in $WI_{[E]}(w)$



when finding a race, compute what the reversed race w looks like


ODPOR: better sets

Weak Initials


$p \in WI_{[E]}(w)$ iff there exists v, w' such that



is equivalent to




To explore w from E , we can start by any process in $WI_{[E]}(w)$



ensure some $p \in WI_{[E]}(w)$ is in the **source set**

ODPOR: wakeup trees

To avoid **sleep set** blocked executions, ODPOR stores trees




A wakeup tree containing sequences P_1P_3 , $P_1P_2P_2$ and $P_2P_2P_3$

Insertion ensures that:

- the exploration of $P_1P_2P_2$ will not be blocked by $\{P_3\}$
- the exploration of $P_2P_2P_3$ will not be blocked by $\{P_1\}$


Limits of this approach



What if the only bug is far from the first guess?


Contribution: Best First (O)DPOR

- Multiple opened states ○
- Corresponding to partially explored executions



- Keeps the optimality from ODPOR
- Allows more freedom in exploration order

BeFS ODPOR differences




- The explored tree is saved as a wakeup tree
- Sleep sets are kept ordered and are updated at the right time
- A procedure garbage collects states when there are no longer needed


Experimental results

MPI example slightly modified

| P_1 | P_2 | P_3 |
|---------------|---------------|-------------------|
| Send(P_3) | Send(P_3) | MPI_Barrier() |
| MPI_Barrier() | MPI_Barrier() | Recv(Any) |
| | | Recv(From P_1) |





In what order?



Exploration strategy ■ DFS ■ Uniform-DFS ■ Uniform-BeFS Branch ■ Uniform-BeFS Step

What's next? - Parallelized exploration



- One **producer** handling the tree
- Multiple **consumers** picking up 
- Distinct explorations happening in parallel

What's next? - Exploration heuristics

Maximize dissimilarity using Fidge-Mattern vector clocks

- Each process stores a clock for each process ($VC \in \mathbb{N}^P$)
- $VC_i[i]$ updates when i does something
- $VC_i[j]$ updates when i and j synchronize over an operation
- States are abstracted as points in $(\mathbb{N}^P)^P$
- Use distance in that space to maximize dissimilarity

What's next? - Exploration heuristics

Maximize dissimilarity using Fidge-Mattern vector clocks

- Each process stores a clock for each process ($VC \in \mathbb{N}^P$)
- $VC_i[i]$ updates when i does something
- $VC_i[j]$ updates when i and j synchronize over an operation
- States are abstracted as points in $(\mathbb{N}^P)^P$
- Use distance in that space to maximize dissimilarity

Using incremental dependencies

- Most bugs only concern 2 or 3 actors (e.g. A and B)
- Start by over-reducing the system as if only A and B were dependent
- Slowly increase the dependencies and repeat

Working on explainability: Why?


```

*****
*** DEADLOCK DETECTED ***
*****
1 actor is still active, awaiting something. Here is its status:
- pid 3 (2@node-10.simgrid.org) simcall CommWait(comm_id:20 src:-1 dst:3 mbox:SMPI-3(id:3))
Counter-example execution trace:
Actor 2 in :0:() ==> simcall: iSend(mbox=3)
Actor 1 in :0:() ==> simcall: iSend(mbox=3)
Actor 1 in :0:() ==> simcall: iRecv(mbox=4)
Actor 1 in :0:() ==> simcall: iRecv(mbox=4)
Actor 2 in :0:() ==> simcall: iSend(mbox=4)
Actor 1 in :0:() ==> simcall: WaitComm(from 2 to 1, mbox=4, no timeout)
Actor 2 in :0:() ==> simcall: iRecv(mbox=5)
Actor 3 in :0:() ==> simcall: iSend(mbox=4)
Actor 1 in :0:() ==> simcall: WaitComm(from 3 to 1, mbox=4, no timeout)
Actor 1 in :0:() ==> simcall: iSend(mbox=5)
Actor 1 in :0:() ==> simcall: iSend(mbox=3)
Actor 1 in :0:() ==> simcall: iRecv(mbox=4)
Actor 1 in :0:() ==> simcall: iRecv(mbox=4)
Actor 2 in :0:() ==> simcall: WaitComm(from 1 to 2, mbox=5, no timeout)
Actor 2 in :0:() ==> simcall: iSend(mbox=4)
Actor 1 in :0:() ==> simcall: WaitComm(from 2 to 1, mbox=4, no timeout)
Actor 2 in :0:() ==> simcall: iRecv(mbox=5)
Actor 3 in :0:() ==> simcall: iRecv(mbox=3)
Actor 3 in :0:() ==> simcall: WaitComm(from 1 to 3, mbox=3, no timeout)
Actor 3 in :0:() ==> simcall: iSend(mbox=4)
Actor 1 in :0:() ==> simcall: WaitComm(from 3 to 1, mbox=4, no timeout)
Actor 1 in :0:() ==> simcall: iSend(mbox=5)
Actor 1 in :0:() ==> simcall: iSend(mbox=3)
Actor 2 in :0:() ==> simcall: WaitComm(from 1 to 2, mbox=5, no timeout)
Actor 3 in :0:() ==> simcall: iRecv(mbox=3)
Actor 3 in :0:() ==> simcall: WaitComm(from 1 to 3, mbox=3, no timeout)
Actor 3 in :0:() ==> simcall: iRecv(mbox=3)
Actor 3 in :0:() ==> simcall: WaitComm(from 2 to 3, mbox=3, no timeout)
Actor 3 in :0:() ==> simcall: iRecv(mbox=3)

```

Mc SimGrid output on a simple example with only two `MPI_Barrier()`.


The critical transition



Critical transition

Let E be an **incorrect** execution,


The critical transition



Critical transition

Let E be an **incorrect** execution,
the **critical transition** is the unique
 $t = (s, a, s') \in E$ s.t.

The critical transition




Critical transition

Let E be an **incorrect** execution, the **critical transition** is the unique $t = (s, a, s') \in E$ s.t.

- every execution from s' is **incorrect**

The critical transition



Critical transition


Let E be an **incorrect** execution, the **critical transition** is the unique $t = (s, a, s') \in E$ s.t.

- every execution from s' is **incorrect**
- there exists a **correct** execution from s

Critical transition: how to compute?

Use reduction and take a decision for the non-explored transitions

- s_{k+1} violates the property
- c_1 is the root of a correct subtree
- Hence, the critical transition is in $\{b_1, \dots, b_{k+1}\}$



Critical transition: what are we missing?

A two process deadlock

| P_1 | P_2 |
|-------------|-------------|
| Lock(a) | Lock(b) |
| Lock(b) | Lock(a) |

- Executions starting by P_1P_2 or P_2P_1 **will** deadlock
- Critical transition is the last executed of $P_1 : \text{Lock}(a)$ and $P_2 : \text{Lock}(b)$
- Possible to retrieve both P_1 and P_2 locks

Critical transition: what are we missing?

A *four* process deadlock

| P_1 | P_2 | P_3 | P_4 |
|-------------|-------------|-------------|-------------|
| Lock(a) | Lock(b) | Lock(c) | Lock(d) |
| Lock(b) | Lock(a) | Lock(d) | Lock(c) |

- Executions starting by ... (24 combinations) **will** deadlock
- Critical transition is the last executed of ... (one of the processes' first action)
- No links between a/b and c/d deadlocks

Conclusion

Contributions

- New reduction algorithms allowing arbitrary search
- Defining and computing critical transitions
- Code integrated within McSimGrid

Future work

- Parallelize the implementation of BeFS ODPOR
- Develop a good benchmark to explore heuristics
- Simplify counter examples using critical sections
- Observe memory access and detect data races with McSimGrid

Time and memory performances

| Benchmark Name | Traces | DPU (UDPOR) | | Nidhugg (ODPOR) | | McSG(BeFS ODPOR) | |
|-------------------|--------|-------------|-------|-----------------|-----|------------------|--------|
| | | Time | Mem | Time | Mem | Time | Mem |
| DISP(5,3) | 1482 | 0.629 | 55M | 6.314 | 65M | 2.080 | 54M |
| DISP(5,4) | 15282 | 6.285 | 135M | 65.034 | 65M | 15.245 | 460M |
| DISP(5,5) | 151032 | 203.785 | 973M | TO | 65M | 154.689 | 4387M |
| DISP(5,6) | | ERR | 1016M | TO | 65M | TO | 17219M |
| MPAT(5) | 3840 | 1.860 | 80M | 1.203 | 64M | 3.927 | 154M |
| MPAT(6) | 46080 | 51.283 | 420M | 16.273 | 64M | 51.426 | 1853M |
| MPAT(7) | 645120 | TO | 1553M | 255.109 | 64M | TO | 19609M |
| MPAT(8) | | TO | 1603M | TO | 64M | TO | 23999M |
| MPC(2,5) | 60 | 0.273 | 51M | 1.038 | 65M | 0.067 | 12M |
| MPC(3,5) | 2958 | 0.937 | 61M | 37.662 | 65M | 2.510 | 81M |
| MPC(4,5) | 313683 | ERR | 63M | TO | 65M | 308.723 | 6684M |
| MPC(5,5) | | TO | 1344M | TO | 65M | TO | 23495M |
| PI(5) | 120 | 0.301 | 43M | ERR | 66M | 0.082 | 11M |
| PI(6) | 720 | 0.468 | 47M | ERR | 66M | 0.441 | 19M |
| PI(7) | 5040 | 1.950 | 66M | ERR | 66M | 3.201 | 77M |
| PI(8) | 40320 | 28.748 | 273M | ERR | 66M | 26.796 | 573M |
| PI(9) | 362880 | TO | 1128M | ERR | 65M | 291.884 | 5291M |
| POKE(7) | 2440 | 1.247 | 84M | 44.736 | 65M | 3.057 | 118M |
| POKE(8) | 3700 | 1.934 | 99M | 146.232 | 65M | 4.913 | 193M |
| POKE(9) | 5332 | 2.913 | 124M | 458.337 | 65M | 7.653 | 302M |
| POKE(10) | 7384 | 4.479 | 152M | TO | 64M | 11.310 | 446M |
| POKE(11) | 9904 | 6.674 | 193M | TO | 65M | 16.247 | 649M |
| POKE(12) | 12940 | 9.969 | 242M | TO | 65M | 22.676 | 910M |
| POKE(13) | 16540 | 14.506 | 310M | ERR | 64M | 30.774 | 1252M |

More results

